

The GPU Rental Cheat Sheet

NightTale · nighttalegames.com

VRAM decides if it works. Bandwidth decides how fast. Disk decides how often you wait.

1 • What you need for what

Job	Minimum	Sweet spot	Speed driver
LLM chat / story text	16GB (8B)	24GB (14B 4-bit)	Memory bandwidth
Image generation (SDXL)	16GB	24-32GB	GPU cores + more cards
LoRA training	24GB	40-80GB	VRAM size
Video generation	24GB	32-80GB	VRAM + architecture
TTS / voice cloning	8GB (CPU ok)	16GB	Small models, cheap

Golden rule: **for batch work, 8 cheap cards beat 1 expensive card.**

2 • The rental lineup

Card	VRAM	Bandwidth	Typical \$/hr	Best for
RTX 3090	24GB	~940 GB/s	\$0.15-0.30	Cheap workers, 14B LLM serving
RTX 4090	24GB	~1,000 GB/s	\$0.30-0.55	Fleet workhorse: images, LoRAs
RTX 5090	32GB	~1,800 GB/s	\$0.50-0.90	32B LLMs, video gen
A100 40/80GB	40/80GB	1.6-2 TB/s	\$0.80-1.60	Training fleets, big models
H100	80GB	~3.3 TB/s	\$2.00-4.00	Serious training only

3 • The 60-second listing check

- **VRAM $\geq$ model + headroom** (weights + KV cache + engine + ~4GB margin).
- **Reliability $\geq$ 98%** — below that, expect the host to die mid-job.
- **Network listed in Mbps** — no number shown = assume it's bad. (Same platform: 60 vs 6,783 Mbps.)
- **Region matches your data** — cross-continent transfers run at ~1 MB/s.
- **RAM $\geq$ 2 $\times$ VRAM, disk $\geq$ dataset + outputs.**
- **Fleet check** — would N cheaper cards beat this one for this job?

4 • Sizing an LLM server

Line item	VRAM
14B model, AWQ 4-bit	~10GB
KV cache (30 players)	~3-5GB
vLLM engine	~2GB
Safety margin	~4GB
Total → 24GB card	~20GB

One 3090 at **\$0.25/hr** + vLLM batching serves **1,000+ active players**.

5 • What it actually costs

Project	Unit cost
30,000 game scenes (8×4090, 5h)	\$0.0009/scene
246 character LoRAs (8×A100, ~15h)	<\$0.50/model
LLM story server (3090, 24/7)	~\$180/mo
Ad video clips (5090, ComfyUI)	\$0.02-0.05/clip

Breakeven on a \$2,500 card: **~6 months of 24/7 rent**. Under → rent. Over → buy.

6 • Operations in four lines

- **flock -n** — one master, ever. No duplicate workers.
- **Skip finished work** — every job resumable; rerunning is always safe.
- **Watchdog via cron every 15 min** — restart what dies.
- **rsync out every 60s** — hosts die; you lose seconds, not hours.

Rent the hardware. Own the models, the pipeline, and the metadata — never the host.

© NightTale — built with this exact playbook: 15 games, 246 models, 30k+ scenes.

[Take the full free course →](#)